

Performers-backed Denoising Diffusion Probabilistic Models

Presentation by Ninniri Matteo (ID: 543873)

October 20, 2022

In this presentation:

We will:

- briefly discuss the theory behind:
 - DDPMs
 - Performers
- Compare our implementation against the original implementation

Table of Contents

1 Theory


- Diffusion models
 - $q(x_t|x_{t-1})$
 - $p_\theta(x_{t-1}|x_t)$
 - Loss
 - Training
 - Sampling
- Performers

Diffusion models

- Latest trend in generative models
- Better results than GANs
- At the core of DALL-E 2 and Imagen

Diffusion process

Gradually add noise to $x_0 \sim q(x_0)$



Given x_{t-1} , x_t is sampled by:

$$q(x_t|x_{t-1}) := \mathcal{N}(x_{t-1}; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I) \quad (1)$$

hence

$$x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}\epsilon \quad (2)$$

where $\epsilon \sim \mathcal{N}(0, I)$ and β_i is a *variance scheduler*

Diffusion process

x_t can be obtained from x_0 directly:


$$x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon \quad (3)$$

where

- $\alpha_t = 1 - \beta_t$
- $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$

Reverse process

Learn to reverse the diffusion



Our target is

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \bar{\mu}(\mathbf{x}_t, t), \bar{\Sigma}(\mathbf{x}_t, t)) \quad (4)$$

which we want to approximate with

$$p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t)) \quad (5)$$

Reverse process

$\bar{\Sigma}(x_t, t)$ is simplified as

$$\bar{\Sigma}(x_t, t) = \sigma_t^2 I \quad (6)$$

where

$$\sigma_t^2 = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t \quad (7)$$

Reverse process

Issue: $q(x_{t-1}|x_t)$ as defined is untractable

\Rightarrow condition it with x_0 : $q(x_{t-1}|x_t, x_0)$:

$$q(x_{t-1}|x_t, x_0) := \mathcal{N}(x_{t-1}; \bar{\mu}(x_t, x_0), \bar{\beta}_t I) \quad (8)$$

Reverse process

Through Bayes' rule (and a lot of steps in between), we obtain

$$\bar{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\alpha_t}}{\beta_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \quad (9)$$

$$\bar{\beta}_t := \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t \quad (10)$$

Reverse process

$$\bar{\mu}_t(x_t, x_0) := \frac{\sqrt{\alpha_t}}{\beta_t} \mathbf{x}_t + \frac{\sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_{t-1}} \mathbf{x}_0 \quad (11)$$

We have x_0 from the diffusion process:

$$\bar{\mu}_t(x_t) := \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_t \right) \quad (12)$$

which can be learned as

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) \quad (13)$$

Reverse process

summarizing:

$$x_{t-1} \sim p_{\theta}(x_{t-1}|x_t) := \mathcal{N}(x_{t-1}; \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}}\epsilon_{\theta}(x_t, t)), \sigma_t I) \quad (14)$$

hence:

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}}\epsilon_{\theta}(x_t, t)) + \sigma_t \epsilon \quad (15)$$

Loss function

We bound on the usual negative log-likelihood:

$$\mathbb{E}[-\log p_{\theta}(x_0)] \tag{16}$$

Loss function

After a lot of intermediate steps:

$$\mathbb{E}_q[-\log p_\theta(x_0)] \leq L := L_T + L_{T-1} + \dots L_1 + L_0 \quad (17)$$

where:

- $L_T = D_{\text{KL}}(q(x_T|x_0)||p(x_T))$
- $L_t = D_{\text{KL}}(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t)); t \in \{2, \dots, T-1\}$
- $L_0 = -\log p_\theta(x_0|x_1)$

L_T

Constant with no trainable parameters
 \Rightarrow No need to optimize it

L_t and L_0

We want to minimize

$$\mathbb{E}_{x_0, \epsilon} \left[\frac{1}{2\sigma_t^2} \|\bar{\mu}_t(x_t) - \mu_\theta(x_t, t)\|^2 \right] \quad (18)$$

We already know both $\bar{\mu}_t(x_t)$ and $\mu_\theta(x_t, t)$:

$$\mathbb{E}_{x_0, \epsilon} \left[\frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_t \right) - \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) \right\|^2 \right] \quad (19)$$

L_t and L_0

Grouping and substituting x_t which we already know

$$\mathbb{E}_{x_0, \epsilon} \left[\frac{(1 - \alpha_t)^2}{2\alpha_t(1 - \bar{\alpha}_t)\sigma_t^2} \|\epsilon - \epsilon_\theta(\sqrt{\alpha_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2 \right] \quad (20)$$

simplifying:

$$L_t := \mathbb{E}_{x_0, \epsilon} [\|\epsilon - \epsilon_\theta(\sqrt{\alpha_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2] \quad (21)$$

\Rightarrow It is all about predicting ϵ_t !

Training

Algorithm 1: Training algorithm

while *True* **do**

$$x_0 \sim q(x_0)$$

$$t \sim \text{Uniform}(\{1, \dots, T\})$$

$$\epsilon \sim \mathcal{N}(0, I)$$

Take gradient descent step on

$$\nabla_{\theta} ||\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)||^2$$

end

Sampling

Algorithm 2: Sampling algorithm

while *True* **do**

$x_T \sim \mathcal{N}(0, I)$

for $t = T, \dots, 1$ **do**

$\epsilon \sim \mathcal{N}(0, I)$ if $t > 1$ else $z = 0$

$x_{t-1} = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}}\epsilon_\theta(x_t, t)) + \sigma_t\epsilon$

end

 return x_0

end

Performers

Standard attention:

$$Att_{\leftrightarrow}(Q, K, V) = D^{-1}AV \quad (22)$$

where

- $A = \exp(\frac{QK^T}{\sqrt{d}})$
- $D = \text{diag}(A1_L)$
- 1_L is the vector of L elements all set to one

Cost: $\mathcal{O}(L^2d)$

Performers

Most kernels can be approximated with

$$K(x, y) = \mathbb{E}[\phi(x)^T \phi(y)] \quad (23)$$

where ϕ is:

$$\phi(x) = \frac{h(x)}{\sqrt{m}} (f_1(x\omega_1^T), \dots, f_1(x\omega_m^T), \dots, f_l(x\omega_1^T), \dots, f_l(x\omega_m^T)) \quad (24)$$

Performers

$$\phi(x) = \frac{h(x)}{\sqrt{m}}(f_1(x\omega_1^T), \dots, f_1(x\omega_m^T), \dots, f_l(x\omega_1^T), \dots, f_l(x\omega_m^T))$$

- $h : \mathbb{R}^d \rightarrow \mathbb{R}$
- $f_1 \dots f_l$ are l functions of the form $\mathbb{R} \rightarrow \mathbb{R}$
- $\omega_1 \dots \omega_m \sim \mathcal{D} \in \mathcal{P}(\mathbb{R}^d)$ are m orthogonal vectors.

Performers

softmax's ϕ has form:

$$\phi(x) = \frac{1}{\sqrt{m}} \exp\left(-\frac{\|x\|^2}{2\sqrt{d}}\right) \left(\exp\left(\frac{x}{\sqrt[4]{d}} \omega_1^T\right), \dots, \exp\left(\frac{x}{\sqrt[4]{d}} \omega_1^T\right)\right) \quad (25)$$

Performers

Summarizing:

$$\widehat{Att}_{\leftrightarrow}(Q, K, V) = \widehat{D}^{-1}(Q'((K')^T V)) \quad (26)$$

where:

- $Q' = \phi(Q)$
- $K' = \phi(K)$
- $\widehat{D} = \text{diag}(Q'((K')^T 1_L))$

Cost: $\mathcal{O}(Ld^2 \log d)$ (if $m = d \log d$)